

[REPORT]



■과 목 명: 운영체제 (SWE3004-42)

■담 당 교 수: 신동균 교수님

■제 출 일: 2022년 5월 23일

■학 과: 수학과

■학 번: 2016314786

■성 명: 김호진

Chap 41. FFS

수학과 김호진 (2016314786)

1. Examine the file in.largefile, and then run the simulator with flag -f in.largefile and -L 4. The latter sets the large-file exception to 4 blocks. What will the resulting allocation look like? Run with -c to check.

in.largefile에서는 총 40개의 블록이 있는 파일을 할당하는데, 실행 결과에서 확인할 수 있듯이 각 10개의 그룹은 30개의 블록으로 구성되어 있다. 그렇기 때문에 flag -f in.largefile만을 사용하여 시뮬레이터를 실행하면 앞선 group부터 차례대로 블록을 할당하게 되고, -L 4를 추가로 사용했을 때는 large-file exception을 4개의 블록으로 설정하여 10개의 그룹에 블록을 4개씩 할당하게 된다.

```
borussen@DESKTOP-L834KLC:~/ostep_ch41$ python3 ffs.py -f in.largefile -c

num_groups:      10
inodes_per_group: 10
blocks_per_group: 30

free data blocks: 259 (of 300)
free inodes:      98 (of 100)

spread inodes?    False
spread data?      False
contig alloc:     1

000000000000000000 1111111111 2222222222
01234567890123456789 0123456789 0123456789

group inodes  data
0 /a----- /aaaaaaaaa aaaaaaaaaa aaaaaaaaaa
1 ----- aaaaaaaaaa a----- -----
2 ----- ----- -----
3 ----- ----- -----
4 ----- ----- -----
5 ----- ----- -----
6 ----- ----- -----
7 ----- ----- -----
8 ----- ----- -----
9 ----- ----- -----

borussen@DESKTOP-L834KLC:~/ostep_ch41$ python3 ffs.py -f in.largefile -L 4 -c

num_groups:      10
inodes_per_group: 10
blocks_per_group: 30

free data blocks: 259 (of 300)
free inodes:      98 (of 100)

spread inodes?    False
spread data?      False
contig alloc:     1

000000000000000000 1111111111 2222222222
01234567890123456789 0123456789 0123456789

group inodes  data
0 /a----- /aaaa----- -----
1 ----- aaaa----- -----
2 ----- aaaa----- -----
3 ----- aaaa----- -----
4 ----- aaaa----- -----
5 ----- aaaa----- -----
6 ----- aaaa----- -----
7 ----- aaaa----- -----
8 ----- aaaa----- -----
9 ----- aaaa----- -----
```

2. Now run with -L 30. What do you expect to see? Once again, turn on -c to see if you were right. You can also use -S to see exactly which blocks were allocated to the file /a.

각 그룹은 최대 30개의 데이터 블록을 가지기 때문에 -L 30으로 시뮬레이터를 실행하게 되면 Large-file exception(-L)을 설정하지 않았을 때와 동일한 결과를 얻게 된다.

```
borussen@DESKTOP-L834KLC:~/ostep_ch41$ python3 ffs.py -f in.largefile -L 30 -S -c
num_groups:      10
inodes_per_group: 10
blocks_per_group: 30

free data blocks: 259 (of 300)
free inodes:      98 (of 100)

spread inodes?    False
spread data?      False
contig alloc:     1

00000000000000000000 1111111111 2222222222
01234567890123456789 0123456789 0123456789

group inodes  data
0 /a----- /aaaaaaaa aaaaaaaaa aaaaaaaaa
1 ----- aaaaaaaaa a----- -----
2 ----- ----- ----- -----
3 ----- ----- ----- -----
4 ----- ----- ----- -----
5 ----- ----- ----- -----
6 ----- ----- ----- -----
7 ----- ----- ----- -----
8 ----- ----- ----- -----
9 ----- ----- ----- -----

symbol inode# filename filetype block_addresses
/      0 / directory 0
a      1 /a regular 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24 25 26 27 28 29 30 31 32 33 34 35 36 37 38 39 40
```

3. Now we will compute some statistics about the file. The first is something we call *filespace*, which is the max distance between any two data blocks of the file or between the inode and any data block. Calculate the filespace of /a. Run ffs.py -f in.largefile -L 4 -T -c to see what it is. Do the same with -L 100. What difference do you expect in filespace as the large-file exception parameter changes from low values to high values?

Large-file exception parameter가 커질수록 많은 그룹으로 분산되지 않기 때문에 -L 값이 증가할 수록 파일의 두 데이터 블록 사이의 최대 거리 또는 inode와 데이터 블록 사이의 최대 거리를 나타내는 filespace는 다음과 같이 감소한다.

```
borussen@DESKTOP-L834KLC:~/ostep_ch41$ python3 ffs.py -f in.largefile -L 4 -T -c
```

```
num_groups:      10
inodes_per_group: 10
blocks_per_group: 30
```

```
free data blocks: 259 (of 300)
free inodes:      98 (of 100)
```

```
spread inodes?   False
spread data?     False
contig alloc:    1
```

```
00000000000000000000 1111111111 2222222222
01234567890123456789 0123456789 0123456789
```

group	inodes	data
0	/a-----	/aaaa-----
1	-----	aaaa-----
2	-----	aaaa-----
3	-----	aaaa-----
4	-----	aaaa-----
5	-----	aaaa-----
6	-----	aaaa-----
7	-----	aaaa-----
8	-----	aaaa-----
9	-----	aaaa-----

```
span: files
file:      /a  filespan: 372
          avg  filespan: 372.00
```

```
span: directories
dir:      /  dirspan: 373
          avg  dirspan: 373.00
```

```
borussen@DESKTOP-L834KLC:~/ostep_ch41$ python3 ffs.py -f in.largefile -L 100 -T -c
```

```
num_groups:      10
inodes_per_group: 10
blocks_per_group: 30
```

```
free data blocks: 259 (of 300)
free inodes:      98 (of 100)
```

```
spread inodes?   False
spread data?     False
contig alloc:    1
```

```
00000000000000000000 1111111111 2222222222
01234567890123456789 0123456789 0123456789
```

group	inodes	data
0	/a-----	/aaaaaaaaa aaaaaaaaaa aaaaaaaaaa
1	-----	aaaaaaaaa a-----
2	-----	-----
3	-----	-----
4	-----	-----
5	-----	-----
6	-----	-----
7	-----	-----
8	-----	-----
9	-----	-----

```
span: files
file:      /a  filespan: 59
          avg  filespan: 59.00
```

```
span: directories
dir:      /  dirspan: 60
          avg  dirspan: 60.00
```

4. Now let's look at a new input file, `in.manyfiles`. How do you think the FFS policy will lay these files out across groups? (you can run with `-v` to see what files and directories are created, or just `cat in.manyfiles`). Run the simulator with `-c` to see if you were right.

해당 FFS 정책에서는 동일한 디렉토리에 있는 파일의 inode 및 데이터 블록을 동일한 그룹에 두어, 3개의 그룹에 파일과 디렉토리를 각각 배치한다.

```
borussen@DESKTOP-L834KLC:~/ostep_ch41$ cat in.manyfiles
file /a 2
file /b 2
file /c 2
file /d 2
file /e 2
file /f 2
file /g 2
file /h 2
file /i 2

dir /j
dir /t

file /t/u 3
file /j/l 1
file /t/v 3
file /j/m 1
file /t/w 3
file /j/n 1
file /t/x 3
file /j/o 1
file /t/y 3
file /j/p 1
file /t/z 3
file /j/q 1
file /t/A 3
file /j/r 1
file /t/B 3
file /j/C 3
```

```
borussen@DESKTOP-L834KLC:~/ostep_ch41$ python3 ffs.py -f in.manyfiles -c

num_groups:      10
inodes_per_group: 10
blocks_per_group: 30

free data blocks: 245 (of 300)
free inodes:      72 (of 100)

spread inodes?   False
spread data?     False
contig alloc:    1

00000000000000000000 1111111111 2222222222
01234567890123456789 0123456789 0123456789

group inodes  data
0 /abcdefghi /aabbccdde effgghhii- -----
1 jlmnopqrC- jlmnopqrCC C-----
2 tuvwxyzAB- tuuuvvvwww xxxyyzzzA ABBBB-----
3 -----
4 -----
5 -----
6 -----
7 -----
8 -----
9 -----
```

5. A metric to evaluate FFS is called *dirspan*. This metric calculates the spread of files within a particular directory, specifically the max distance between the inodes and data blocks of all files in the directory and the inode and data block of the directory itself. Run with `in.manyfiles` and the `-T` flag, and calculate the *dirspan* of the three directories. Run with `-c` to check. How good of a job does FFS do in minimizing *dirspan*?

*dirspan*은 inode block과 data block 사이의 span을 구성하는 inode region의 크기에 큰 영향을 받는다. 해당 FFS의 경우, 파일이 특정 그룹을 overflow 하기에 충분한 공간을 차지하지 않기 때문에 *dirspan*을 최소화하는 데 도움이 된다.

```
borussen@DESKTOP-L834KLC:~/ostep_ch41$ python3 ffs.py -f in.manyfiles -T -c

num_groups:      10
inodes_per_group: 10
blocks_per_group: 30

free data blocks: 245 (of 300)
free inodes:      72 (of 100)

spread inodes?    False
spread data?      False
contig alloc:     1

000000000000000000 111111111 222222222
01234567890123456789 0123456789 0123456789

group inodes      data
0 /abcdefghi /aabbccdde effgghhii- -----
1 jlmnopqrC- jlmnopqrCC C-----
2 tuvwxzAB- tuuuvvwww xxxyyzzzA AABBB-----
3 -----
4 -----
5 -----
6 -----
7 -----
8 -----
9 -----

span: files
file:      /a filespan: 11
file:      /b filespan: 12
file:      /c filespan: 13
file:      /d filespan: 14
file:      /e filespan: 15
file:      /f filespan: 16
file:      /g filespan: 17
file:      /h filespan: 18
file:      /i filespan: 19
file:      /t/u filespan: 12
file:      /j/l filespan: 10
file:      /t/v filespan: 14
file:      /j/m filespan: 10
file:      /t/w filespan: 16
file:      /j/n filespan: 10
file:      /t/x filespan: 18
file:      /j/o filespan: 10
file:      /t/y filespan: 20
file:      /j/p filespan: 10
file:      /t/z filespan: 22
file:      /j/q filespan: 10
file:      /t/A filespan: 24
file:      /j/r filespan: 10
file:      /t/B filespan: 26
file:      /j/C filespan: 12
file:      avg filespan: 14.76

span: directories
dir:      / dirspan: 28
dir:      /j dirspan: 20
dir:      /t dirspan: 34
dir:      avg dirspan: 27.33
```

6. Now change the size of the inode table per group to 5 (-I 5). How do you think this will change the layout of the files? Run with -c to see if you were right. How does it affect the dirspan?

그룹당 inode table의 크기를 5(-I 5)로 변경하면, 같은 디렉토리의 파일이 여러 그룹으로 분산되어 dirspan이 증가하게 된다.

```
borussen@DESKTOP-L834KLC:~/ostep_ch41$ python3 ffs.py -f in.manyfiles -I 5 -T -c

num_groups:      10
inodes_per_group: 10
blocks_per_group: 30

free data blocks: 245 (of 300)
free inodes:      72 (of 100)

spread inodes?    True
spread data?      False
contig alloc:     1

      00000000000000000000 1111111111 2222222222
      01234567890123456789 0123456789 0123456789

group inodes  data
  0 /jy----- /jyyy----- -----
  1 atp----- aatp----- -----
  2 buz----- bbuuuzzz-- -----
  3 clq----- ccldq----- -----
  4 dvA----- ddvvvAAA-- -----
  5 emr----- eemr----- -----
  6 fwB----- ffwwwBBB-- -----
  7 gnC----- ggnCCC---- -----
  8 hx----- hxxxx----- -----
  9 io----- iio----- -----

span: files
  file:      /a  filespan: 11
  file:      /b  filespan: 11
  file:      /c  filespan: 11
  file:      /d  filespan: 11
  file:      /e  filespan: 11
  file:      /f  filespan: 11
  file:      /g  filespan: 11
  file:      /h  filespan: 11
  file:      /i  filespan: 11
  file:      /t/u filespan: 13
  file:      /j/l filespan: 11
  file:      /t/v filespan: 13
  file:      /j/m filespan: 11
  file:      /t/w filespan: 13
  file:      /j/n filespan: 11
  file:      /t/x filespan: 13
  file:      /j/o filespan: 11
  file:      /t/y filespan: 12
  file:      /j/p filespan: 11
  file:      /t/z filespan: 15
  file:      /j/q filespan: 11
  file:      /t/A filespan: 15
  file:      /j/r filespan: 11
  file:      /t/B filespan: 15
  file:      /j/C filespan: 13
              avg  filespan: 11.92

span: directories
  dir:      /  dirspan: 371
  dir:      /j dirspan: 371
  dir:      /t dirspan: 332
              avg  dirspan: 358.00
```

7. Which group should FFS place inode of a new directory in? The default (simulator) policy looks for the group with the most free inodes. A different policy looks for a set of groups with the most free inodes. For example, if you run with -A 2, when allocating a new directory, the simulator will look at groups in pairs and pick the best pair for the allocation. Run `./ffs.py -f in.manyfiles -I 5 -A 2 -c` to see how allocation changes with this strategy. How does it affect dirspan? Why might this policy be good?

-A 2로 실행하는 경우, 해당 정책에 의해서 디렉토리와 해당 파일의 inode 사이에 다른 디렉토리가 들어가지 않기 때문에 dirspan이 감소한다.

```
borussen@DESKTOP-L834KLC:~/ostep_ch41$ python3 ffs.py -fin.manyfiles -I 5 -A 2 -T -c

num_groups:      10
inodes_per_group: 10
blocks_per_group: 30

free data blocks: 245 (of 300)
free inodes:      72 (of 100)

spread inodes?    True
spread data?      False
contig alloc:     1

000000000000000000 1111111111 2222222222
01234567890123456789 0123456789 0123456789

group inodes  data
0 /ejmyr---- /eejmyyr- -----
1 -----
2 aftpB---- aafftwwpB BB-----
3 -----
4 bgunzC---- bbgguunzz zCCC-----
5 -----
6 chlXq---- cchhlxxxq- -----
7 -----
8 divoA---- ddiivvvoAA A-----
9 -----

span: files
file:      /a  filespan: 11
file:      /b  filespan: 11
file:      /c  filespan: 11
file:      /d  filespan: 11
file:      /e  filespan: 11
file:      /f  filespan: 12
file:      /g  filespan: 12
file:      /h  filespan: 12
file:      /i  filespan: 12
file:      /t/u filespan: 14
file:      /j/l filespan: 12
file:      /t/v filespan: 14
file:      /j/m filespan: 11
file:      /t/w filespan: 14
file:      /j/n filespan: 14
file:      /t/x filespan: 14
file:      /j/o filespan: 14
file:      /t/y filespan: 13
file:      /j/p filespan: 14
file:      /t/z filespan: 16
file:      /j/q filespan: 14
file:      /t/A filespan: 16
file:      /j/r filespan: 13
file:      /t/B filespan: 16
file:      /j/C filespan: 18
          avg  filespan: 13.20

span: directories
dir:      /  dirspan: 333
dir:      /j dirspan: 335
dir:      /t dirspan: 336
          avg dirspan: 334.67
```


8. One last policy change we will explore relates to file fragmentation. Run `./ffs.py -f in.fragmented -v` and see if you can predict how the files that remain are allocated. Run with `-c` to confirm your answer. What is interesting about the data layout of file `/i`? Why is it problematic?

`./ffs.py -f in.fragmented -v`를 실행했을 때, 파일 `/i`의 data layout은 fragmented 하고 non-contiguous 하다. 이럴 경우, 파일을 read 하거나 write 하기 위해서 disk가 더 많은 검색 및 회전을 해야 하기 때문에 속도가 느려지는 문제가 발생한다.

```
borussen@DESKTOP-L834KLC:~/ostep_ch41$ python3 ffs.py -f in.fragmented -T -v -c
op create /a [size:1] ->success
op create /b [size:1] ->success
op create /c [size:1] ->success
op create /d [size:1] ->success
op create /e [size:1] ->success
op create /f [size:1] ->success
op create /g [size:1] ->success
op create /h [size:1] ->success
op delete /a ->success
op delete /c ->success
op delete /e ->success
op delete /g ->success
op create /i [size:8] ->success

num_groups:      10
inodes_per_group: 10
blocks_per_group: 30

free data blocks: 287 (of 300)
free inodes:      94 (of 100)

spread inodes?    False
spread data?      False
contig alloc:     1

00000000000000000000 1111111111 2222222222
01234567890123456789 0123456789 0123456789

group inodes  data
0 /ib-d-f-h- /ibidifihi iii-----
1 -----
2 -----
3 -----
4 -----
5 -----
6 -----
7 -----
8 -----
9 -----

span: files
file:      /b  filespan: 10
file:      /d  filespan: 10
file:      /f  filespan: 10
file:      /h  filespan: 10
file:      /i  filespan: 21
           avg  filespan: 12.20

span: directories
dir:      /  dirspan: 22
           avg  dirspan: 22.00
```

9. A new policy, which we call *contiguous allocation* (-C), tries to ensure that each file is allocated contiguously. Specifically, with -C n, the file system tries to ensure that n contiguous blocks are free within a group before allocating a block. Run ./ffs.py -f in.fragmented -v -C 2 -c to see the difference. How does layout change as the parameter passed to -C increases? Finally, how does -C affect filespan and dirspan?

./ffs.py -f in.fragmented -v -C 2 -c를 실행하면, 파일 /i의 data layout이 contiguous하게 할당된다. 만약 -C로 전달된 파라미터가 증가한다면 -C보다 작은 파일이 데이터 블록 영역에서 external fragmentation을 유발하게 된다. 그러므로 -C는 filespan과 dirspan을 증가시킨다.

```
borussen@DESKTOP-L834KLC:~/ostep_ch41$ python3 ffs.py -f in.fragmented -T -v -C 2 -c
op create /a [size:1] ->success
op create /b [size:1] ->success
op create /c [size:1] ->success
op create /d [size:1] ->success
op create /e [size:1] ->success
op create /f [size:1] ->success
op create /g [size:1] ->success
op create /h [size:1] ->success
op delete /a ->success
op delete /c ->success
op delete /e ->success
op delete /g ->success
op create /i [size:8] ->success

num_groups:      10
inodes_per_group: 10
blocks_per_group: 30

free data blocks: 287 (of 300)
free inodes:      94 (of 100)

spread inodes?    False
spread data?      False
contig alloc:     2

00000000000000000000 111111111 222222222
01234567890123456789 0123456789 0123456789

group inodes  data
0 /ib-d-f-h- /-b-d-f-hi iiiiii--- -----
1 -----
2 -----
3 -----
4 -----
5 -----
6 -----
7 -----
8 -----
9 -----

span: files
file:      /b  filespan: 10
file:      /d  filespan: 10
file:      /f  filespan: 10
file:      /h  filespan: 10
file:      /i  filespan: 25
          avg  filespan: 13.00

span: directories
dir:      /  dirspan: 26
          avg  dirspan: 26.00
```